

Introduction_to_Relational_Databases

Pop quiz: The relational model

Which of the following is not part of the relational model?

Options:

1. Each row or record in a table represents an instance of an entity type.
2. Each column in a table represents an attribute or feature of an instance.
3. Every table contains a primary key column, which has a unique entry for each row.
4. A database consists of at least 3 tables.
5. There are relations between tables.

Answer: 4

1.1 Creating a database engine in Python

Here, you're going to fire up your very first SQL engine. You'll create an engine to connect to the SQLite database 'Chinook.sqlite', which is in your working directory. Remember that to create an engine to connect to 'Northwind.sqlite', Hugo executed the command

```
engine = create_engine('sqlite:///Northwind.sqlite')
```

Here, 'sqlite:///Northwind.sqlite' is called the connection string to the SQLite database Northwind.sqlite. A little bit of background on the [Chinook database \(https://github.com/lerocha/chinook-database\)](https://github.com/lerocha/chinook-database): the Chinook database contains information about a semi-fictional digital media store in which media data is real and customer, employee and sales data has been manually created.

Why the name Chinook, you ask? According to their [website \(https://github.com/lerocha/chinook-database\)](https://github.com/lerocha/chinook-database),

The name of this sample database was based on the Northwind database. Chinooks are winds in the interior West of North America, where the Canadian Prairies and Great Plains meet various mountain ranges. Chinooks are most prevalent over southern Alberta in Canada. Chinook is a good name choice for a database that intends to be an alternative to Northwind.

Instructions

- Import the function `create_engine` from the module `sqlalchemy`.
- Create an engine to connect to the SQLite database 'Chinook.sqlite' and assign it to `engine`.

In [1]:

```
1 # Import necessary module
2 from sqlalchemy import create_engine
3
4 # Create engine: engine
5 engine = create_engine('sqlite:///Chinook.sqlite')
```



1.2 What are the tables in the database?

In this exercise, you'll once again create an engine to connect to `'Chinook.sqlite'`. Before you can get any data out of the database, however, you'll need to know what tables it contains!

To this end, you'll save the table names to a list using the method `table_names()` on the engine and then you will print the list.

Instructions:

- Import the function `create_engine` from the module `sqlalchemy`.
- Create an engine to connect to the SQLite database `'Chinook.sqlite'` and assign it to `engine`.
- Using the method `table_names()` on the engine `engine`, assign the table names of `'Chinook.sqlite'` to the variable `table_names`.
- Print the object `table_names` to the shell.

In [2]:



```
1 # Import necessary module
2 from sqlalchemy import create_engine
3
4 # Create engine: engine
5 engine = create_engine('sqlite:///Chinook.sqlite')
6
7 # Save the table names to a list:
8
9 table_names = engine.table_names()
10 # Print the table names to the shell
11 print(table_names)
```

```
['Album', 'Artist', 'Customer', 'Employee', 'Genre', 'Invoice', 'InvoiceLine', 'MediaType', 'Playlist', 'PlaylistTrack', 'Track']
```

2. Querying relational databases in Python

2.1 The Hello World of SQL Queries!

Now, it's time for liftoff! In this exercise, you'll perform the Hello World of SQL queries, `SELECT`, in order to retrieve all columns of the table `Album` in the Chinook database. Recall that the query `SELECT *` selects all columns.

Instructions

- Open the engine connection as `con` using the method `connect()` on the engine.
- Execute the query that selects ALL columns from the `Album` table. Store the results in `rs`.
- Store all of your query results in the DataFrame `df` by applying the `fetchall()` method to the results `rs`.
- Close the connection!

In [3]:



```
1 # Import packages
2 from sqlalchemy import create_engine
3 import pandas as pd
4
5 # Create engine: engine
6 engine = create_engine('sqlite:///Chinook.sqlite')
7
8 # Open engine connection:
9 con = engine.connect()
10
11 # Perform query: rs
12 rs = con.execute('SELECT * FROM Album')
13
14 # Save results of the query to DataFrame: df
15 df = pd.DataFrame(rs.fetchall())
16
17 # Close connection
18 con.close()
19
20 # Print head of DataFrame df
21 print(df.head())
```

	0	1	2
0	1	For Those About To Rock We Salute You	1
1	2	Balls to the Wall	2
2	3	Restless and Wild	2
3	4	Let There Be Rock	1
4	5	Big Ones	3

2.2 Customizing the Hello World of SQL Queries

Congratulations on executing your first SQL query! Now you're going to figure out how to customize your query in order to:

- Select specified columns from a table;
- Select a specified number of rows;
- Import column names from the database table.
- Recall that Hugo performed a very similar query customization in the video:

```
engine = create_engine('sqlite:///Northwind.sqlite')
```

```
with engine.connect() as con: rs = con.execute("SELECT OrderID, OrderDate, ShipName
FROM Orders") df = pd.DataFrame(rs.fetchmany(size=5)) df.columns = rs.keys()
```

Packages have already been imported as follows:

```
from sqlalchemy import create_engine import pandas as pd The engine has also already been
created:
```

```
engine = create_engine('sqlite:///Chinook.sqlite')
```

The engine connection is already open with the statement

```
with engine.connect() as con:
```

All the code you need to complete is within this context.

Instructions

- Execute the SQL query that selects the columns `LastName` and `Title` from the `Employee` table. Store the results in the variable `rs`.
- Apply the method `fetchmany()` to `rs` in order to retrieve 3 of the records. Store them in the DataFrame `df`.
- Using the `rs` object, set the DataFrame's column names to the corresponding names of the table columns.

In [4]:



```
1 # Open engine in context manager
2 # Perform query and save results to DataFrame: df
3 with engine.connect() as con:
4     rs = con.execute("SELECT LastName, Title FROM Employee")
5     df = pd.DataFrame(rs.fetchmany(size=3))
6     df.columns = rs.keys()
7
8 # Print the length of the DataFrame df
9 print(len(df))
10
11 # Print the head of the DataFrame df
12 print(df.head())
```

```
3
  LastName      Title
0   Adams  General Manager
1  Edwards   Sales Manager
2  Peacock Sales Support Agent
```

3.2 Filtering your database records using SQL's WHERE

You can now execute a basic SQL query to select records from any table in your database and you can also perform simple query customizations to select particular columns and numbers of rows.

There are a couple more standard SQL query chops that will aid you in your journey to becoming an SQL ninja.

Let's say, for example that you wanted to get all records from the `Customer` table of the `Chinook` database for which the `Country` is `'Canada'`. You can do this very easily in SQL using a `SELECT` statement followed by a `WHERE` clause as follows:

```
SELECT * FROM Customer WHERE Country = 'Canada'
```

In fact, you can filter any `SELECT` statement by any condition using a `WHERE` clause. This is called filtering your records.

Packages are already imported as follows:

```
import pandas as pd from sqlalchemy import create_engine Query away!
```

Instructions

- Complete the argument of `create_engine()` so that the engine for the SQLite database 'Chinook.sqlite' is created.
- Execute the query that selects all records from the `Employee` table where 'EmployeeId' is greater than or equal to 6. Use the `>=` operator and assign the results to `rs`.
- Apply the method `fetchall()` to `rs` in order to fetch all records in `rs`. Store them in the DataFrame `df`.
- Using the `rs` object, set the DataFrame's column names to the corresponding names of the table columns.

In [5]:



```
1 # Create engine: engine
2 engine = create_engine('sqlite:///Chinook.sqlite')
3
4 # Open engine in context manager
5 # Perform query and save results to DataFrame: df
6 with engine.connect() as con:
7     rs = con.execute('SELECT * FROM Employee WHERE EmployeeId>= 6')
8     df = pd.DataFrame(rs.fetchall())
9     df.columns = rs.keys()
10
11 # Print the head of the DataFrame df
12 print(df.head())
```

	EmployeeId	LastName	FirstName	Title	ReportsTo	BirthDat
e \						
0	6	Mitchell	Michael	IT Manager	1	1973-07-01 00:00:00
0						
1	7	King	Robert	IT Staff	6	1970-05-29 00:00:00
0						
2	8	Callahan	Laura	IT Staff	6	1968-01-09 00:00:00
0						

	HireDate	Address	City	State	Countr
y \					
0	2003-10-17 00:00:00	5827 Bowness Road NW	Calgary	AB	Canad
a					
1	2004-01-02 00:00:00	590 Columbia Boulevard West	Lethbridge	AB	Canad
a					
2	2004-03-04 00:00:00	923 7 ST NW	Lethbridge	AB	Canad
a					

	PostalCode	Phone	Fax	Email
0	T3B 0C5	+1 (403) 246-9887	+1 (403) 246-9899	michael@chinookcorp.com
1	T1K 5N8	+1 (403) 456-9986	+1 (403) 456-8485	robert@chinookcorp.com
2	T1H 1Y8	+1 (403) 467-3351	+1 (403) 467-8772	laura@chinookcorp.com

In [6]:



```
1 with engine.connect() as con:
2     rs = con.execute('SELECT * FROM Employee WHERE EmployeeId>= 6' )
3     print(rs.keys())
```

```
['EmployeeId', 'LastName', 'FirstName', 'Title', 'ReportsTo', 'BirthDate',
'HireDate', 'Address', 'City', 'State', 'Country', 'PostalCode', 'Phone', 'Fax', 'Email']
```

3.3 Ordering your SQL records with ORDER BY

You can also order your SQL query results. For example, if you wanted to get all records from the Customer table of the Chinook database and order them in increasing order by the column SupportRepId, you could do so with the following query:

```
"SELECT * FROM Customer ORDER BY SupportRepId"
```

In fact, you can order any `SELECT` statement by any column.

Packages are already imported as follows:

```
import pandas as pd from sqlalchemy import create_engine
```

Get querying!

Instructions

- Using the function `create_engine()`, create an engine for the SQLite database `Chinook.sqlite` and assign it to the variable `engine`.
- In the context manager, execute the query that selects all records from the `Employee` table and orders them in increasing order by the column `BirthDate`. Assign the result to `rs`.
- In a call to `pd.DataFrame()`, apply the method `fetchall()` to `rs` in order to fetch all records in `rs`. Store them in the DataFrame `df`.
- Set the DataFrame's column names to the corresponding names of the table columns.

In [7]:



```
1 # Create engine: engine
2
3 engine = create_engine('sqlite:///Chinook.sqlite')
4 # Open engine in context manager
5 with engine.connect() as con:
6     rs = con.execute("SELECT * From Employee ORDER BY BirthDate ")
7     df = pd.DataFrame(rs.fetchall())
8
9     # Set the DataFrame's column names
10
11     df.columns = rs.keys()
12 # Print head of DataFrame
13 print(df.head())
```

	EmployeeId	LastName	FirstName	Title	ReportsTo	\
0	4	Park	Margaret	Sales Support Agent	2.0	
1	2	Edwards	Nancy	Sales Manager	1.0	
2	1	Adams	Andrew	General Manager	NaN	
3	5	Johnson	Steve	Sales Support Agent	2.0	
4	8	Callahan	Laura	IT Staff	6.0	

	BirthDate	HireDate	Address	City
0	1947-09-19 00:00:00	2003-05-03 00:00:00	683 10 Street SW	Calgary
1	1958-12-08 00:00:00	2002-05-01 00:00:00	825 8 Ave SW	Calgary
2	1962-02-18 00:00:00	2002-08-14 00:00:00	11120 Jasper Ave NW	Edmonton
3	1965-03-03 00:00:00	2003-10-17 00:00:00	7727B 41 Ave	Calgary
4	1968-01-09 00:00:00	2004-03-04 00:00:00	923 7 ST NW	Lethbridge

	State	Country	PostalCode	Phone	Fax	\
0	AB	Canada	T2P 5G3	+1 (403) 263-4423	+1 (403) 263-4289	
1	AB	Canada	T2P 2T3	+1 (403) 262-3443	+1 (403) 262-3322	
2	AB	Canada	T5K 2N1	+1 (780) 428-9482	+1 (780) 428-3457	
3	AB	Canada	T3B 1Y7	1 (780) 836-9987	1 (780) 836-9543	
4	AB	Canada	T1H 1Y8	+1 (403) 467-3351	+1 (403) 467-8772	

	Email
0	margaret@chinookcorp.com
1	nancy@chinookcorp.com
2	andrew@chinookcorp.com
3	steve@chinookcorp.com
4	laura@chinookcorp.com

4. Querying relational databases directly with pandas

4.1 Pandas and The Hello World of SQL Queries!

Here, you'll take advantage of the power of `pandas` to write the results of your SQL query to a `DataFrame` in one swift line of Python code!

You'll first import `pandas` and create the SQLite `'Chinook.sqlite'` engine. Then you'll query the database to select all records from the `Album` table.

Recall that to select all records from the `Orders` table in the `Northwind` database, Hugo executed the following command:

```
df = pd.read_sql_query("SELECT * FROM Orders", engine)
```

Instructions

- Import the `pandas` package using the alias `pd`.
- Using the function `create_engine()`, create an engine for the SQLite database `Chinook.sqlite` and assign it to the variable `engine`.
- Use the `pandas` function `read_sql_query()` to assign to the variable `df` the `DataFrame` of results from the following query: select all records from the table `Album`.

In [8]:



```
1 # Import packages
2 from sqlalchemy import create_engine
3 import pandas as pd
4
5 # Create engine:
6 engine = create_engine("sqlite:///Chinook.sqlite")
7
8 # Execute query and store records in DataFrame: df
9 df = pd.read_sql_query("SELECT * FROM Album", engine)
10
11 # Print head of DataFrame
12 print(df.head())
13
14 # Open engine in context manager and store query result in df1
15 with engine.connect() as con:
16     rs = con.execute("SELECT * FROM Album")
17     df1 = pd.DataFrame(rs.fetchall())
18     df1.columns = rs.keys()
19
20 # Confirm that both methods yield the same result
21 print(df.equals(df1))
```

	AlbumId	Title	ArtistId
0	1	For Those About To Rock We Salute You	1
1	2	Balls to the Wall	2
2	3	Restless and Wild	2
3	4	Let There Be Rock	1
4	5	Big Ones	3

True

Pandas for more complex querying

Here, you'll become more familiar with the `pandas` function `read_sql_query()` by using it to execute a more complex query: a `SELECT` statement followed by both a `WHERE` clause AND an `ORDER BY` clause.

You'll build a DataFrame that contains the rows of the Employee table for which the EmployeeId is greater than or equal to 6 and you'll order these entries by BirthDate .

Instructions

- Using the function create_engine(), create an engine for the SQLite database Chinook.sqlite and assign it to the variable engine.
- Use the pandas function read_sql_query() to assign to the variable df the DataFrame of results from the following query: **select** all records from the Employee table **where** the EmployeeId is greater than or equal to 6 and **ordered by** BirthDate (make sure to use WHERE and ORDER BY in this precise order).

In [9]:

```
1 # Import packages
2 from sqlalchemy import create_engine
3 import pandas as pd
4
5 # Create engine: engine
6
7 engine = create_engine("sqlite:///Chinook.sqlite")
8 # Execute query and store records in DataFrame: df
9
10 df=pd.read_sql_query("SELECT * FROM Employee WHERE EmployeeId>=6 ORDER BY BirthDate",engine)
11
12 # Print head of DataFrame
13 print(df.head())
14
```

EmployeeId	LastName	FirstName	Title	ReportsTo	BirthDate	
8	Callahan	Laura	IT Staff	6	1968-01-09 00:00:00	
7	King	Robert	IT Staff	6	1970-05-29 00:00:00	
6	Mitchell	Michael	IT Manager	1	1973-07-01 00:00:00	
HireDate		Address		City	State	Country
2004-03-04 00:00:00	923 7 ST NW		Lethbridge	AB	Canada	
2004-01-02 00:00:00	590 Columbia Boulevard West		Lethbridge	AB	Canada	
2003-10-17 00:00:00	5827 Bowness Road NW		Calgary	AB	Canada	
PostalCode		Phone		Fax	Email	
T1H 1Y8	+1 (403) 467-3351	+1 (403) 467-8772	laura@chinookcorp.com			
T1K 5N8	+1 (403) 456-9986	+1 (403) 456-8485	robert@chinookcorp.com			
T3B 0C5	+1 (403) 246-9887	+1 (403) 246-9899	michael@chinookcorp.com			

In [10]:



```
1 # Import packages
2 from sqlalchemy import create_engine
3 import pandas as pd
4
5 # Create engine: engine
6
7 engine = create_engine("sqlite:///Chinook.sqlite")
8 # Execute query and store records in DataFrame: df
9
10 df=pd.read_sql_query("SELECT * From Artist",engine)
11 df1=pd.read_sql_query("SELECT * From Album",engine)
12 # Print head of DataFrame
13 print(df.head())
14 print(df1.head())
```

	ArtistId	Name
0	1	AC/DC
1	2	Accept
2	3	Aerosmith
3	4	Alanis Morissette
4	5	Alice In Chains

	AlbumId	Title	ArtistId
0	1	For Those About To Rock We Salute You	1
1	2	Balls to the Wall	2
2	3	Restless and Wild	2
3	4	Let There Be Rock	1
4	5	Big Ones	3

The power of SQL lies in relationships between tables: INNER JOIN

Here, you'll perform your first INNER JOIN! You'll be working with your favourite SQLite database, Chinook.sqlite. For each record in the Album table, you'll extract the Title along with the Name of the Artist. The latter will come from the Artist table and so you will need to INNER JOIN these two tables on the ArtistID column of both.

Recall that to INNER JOIN the Orders and Customers tables from the Northwind database, Hugo executed the following SQL query:

```
"SELECT OrderID, CompanyName FROM Orders INNER JOIN Customers on
Orders.CustomerID = Customers.CustomerID"
```

The following code has already been executed to import the necessary packages and to create the engine:

```
import pandas as pd from sqlalchemy import create_engine engine =
create_engine('sqlite:///Chinook.sqlite')
```

Instructions

- Assign to rs the results from the following query: select all the records, extracting the Title of the record and Name of the artist of each record from the Album table and the Artist table, respectively. To do so, INNER JOIN these two tables on the ArtistID column of both.

- In a call to `pd.DataFrame()` , apply the method `fetchall()` to `rs` in order to fetch all records in `rs` . Store them in the DataFrame `df` .
- Set the DataFrame's column names to the corresponding names of the table columns.

In [11]:



```
1 # Open engine in context manager
2 engine = create_engine("sqlite:///Chinook.sqlite")
3
4 # Perform query and save results to DataFrame: df
5 with engine.connect() as con:
6     rs = con.execute("SELECT Title, Name FROM Album INNER JOIN Artist on Album.ArtistId
7     df = pd.DataFrame(rs.fetchall())
8     df.columns = rs.keys()
9     df.head()
10
11 # Print head of DataFrame df
12 print(df.head())
13
```

	Title	Name
0	For Those About To Rock We Salute You	AC/DC
1	Balls to the Wall	Accept
2	Restless and Wild	Accept
3	Let There Be Rock	AC/DC
4	Big Ones	Aerosmith

Filtering your INNER JOIN

Congrats on performing your first INNER JOIN! You're now going to finish this chapter with one final exercise in which you perform an INNER JOIN and filter the result using a WHERE clause.

Recall that to INNER JOIN the Orders and Customers tables from the Northwind database, Hugo executed the following SQL query:

"SELECT OrderID, CompanyName FROM Orders INNER JOIN Customers on Orders.CustomerID = Customers.CustomerID" The following code has already been executed to import the necessary packages and to create the engine:

```
import pandas as pd from sqlalchemy import create_engine engine = create_engine('sqlite:///Chinook.sqlite')
Instructions 100 XP Use the pandas function read_sql_query() to assign to the variable df the DataFrame of
results from the following query: select all records from PlaylistTrack INNER JOIN Track on
PlaylistTrack.TrackId = Track.TrackId that satisfy the condition Milliseconds < 250000.
```

In [12]:



```
1 # Execute query and store records in DataFrame: df
2
3 df = pd.read_sql_query('SELECT * FROM PlaylistTrack INNER JOIN Track on PlaylistTrack.TrackId = Track.TrackId')
4
5 # Print head of DataFrame
6 print(df.head())
```

	PlaylistId	TrackId	TrackId	Name	AlbumId	MediaTypeId	\
0	1	3390	3390	One and the Same	271	2	
1	1	3392	3392	Until We Fall	271	2	
2	1	3393	3393	Original Fire	271	2	
3	1	3394	3394	Broken City	271	2	
4	1	3395	3395	Somedays	271	2	

	GenreId	Composer	Milliseconds	Bytes	UnitPrice
0	23	None	217732	3559040	0.99
1	23	None	230758	3766605	0.99
2	23	None	218916	3577821	0.99
3	23	None	228366	3728955	0.99
4	23	None	213831	3497176	0.99